Automatic Non-functional Testing of Code Generators Families

Mohamed Boussaa   Olivier Barais
Benoit Baudry

Diverse team INRIA, Rennes, France

{mohamed.boussaa, olivier.barais,
benoit.baudry}@inria.fr

Gerson Sunyé

AtlanMod team INRIA, Nantes, France
gerson.sunye@inria.fr

Abstract

The intensive use of generative programming techniques provides an elegant engineering solution to deal with the heterogeneity of platforms or technological stacks. The use of Domain Specifics Language, for example, leads to the creation of numerous code generators and compilers that will automatically translate high-level system specifications into multi-target executable code. Producing correct and efficient code generator is complex and error-prone. Although software designers provide generally high-level test suites to verify the functional outcome of generated code, it remains challenging and tedious to verify the behavior of produced code in terms of non-functional properties.

This paper describes a practical approach based on a runtime monitoring infrastructure to automatically check potential inefficient code generator. This infrastructure, based on system containers as execution platforms, allows code-generator developers to evaluate the consistency and coherence of generated code regarding the non-functional properties. This approach provides a fine-grained understanding of resource consumption and analysis of components behavior. We evaluate our approach by analyzing the non-functional properties of HAXE, a popular high-level programming language that involves a set of cross-platform code generators able to compile to different target platforms. Our experimental results show that our approach is able to detect some non-functional inconsistencies within HAXE code generators.

Categories and Subject Descriptors   D.3.4 [*Programming Languages*]: Processors compilers, code generation

General Terms   Generative Programming, Testing, Components

Keywords   code quality, non-functional properties, code generator, testing

1. Introduction

Nowadays, the intensive use of generative programming techniques has become a common practice for software development to tame the runtime platforms heterogeneity that exists in several domains such as mobile or Internet of Things development. Generative programming techniques reduce the development and maintenance effort by developing at a higher-level of abstraction through the use of domain-specific languages [1] (DSLs). A code generator can be used to transform source code programs/models represented in a graphical/textual modeling language to general purpose programming languages such as C, Java, C++, PHP, JavaScript etc. In turn, generated code is transformed into machine code (binaries) using a set of specific compilers.

However, code generators can be difficult to understand since they involve a set of complex and heterogeneous technologies which make the task of performing design, implementation, and testing very hard and time-consuming [4, 6]. Code generators have to respect different requirements which preserve software reliability and quality. In fact, faulty code generators can generate defective software artifacts which range from uncompilable or semantically dysfunctional code that causes serious damage to the target platform, to non-functional bugs which lead to poor-quality code that can affect system reliability and performance (e.g., high resource usage, high execution time, etc.).

In order to check the correctness of the code generation process, developers often use to define (at design time) a set of test cases that verify the functional outcome of generated code. After code generation, test suites are executed within each target platform. This may lead to either a correct behavior (i.e., expected output) or a failure (i.e., crashes, bugs). On the other hand, for performance testing of code generators, developers need to deploy and execute software artifacts within different execution platforms. Then, they have to collect and compare information about the performance and efficiency of the generated code. Finally, they report issues related to the code generation process such as incorrect typing, memory management leaks, etc. Currently there is no automatic solution to check the performance issues such as huge memory/CPU consumption of the generated code. In fact, developers use manually several platform-specific pro-

filers, trackers, and monitoring tools [2, 5] in order to find some inconsistencies or bugs during code execution. Ensuring the code quality of generated code can refer to several non-functional properties such as code size, resource or energy consumption, execution time, among others [12]. Testing the non-functional properties of code generators remains a challenging and time-consuming task because developers have to analyze and verify code for each target platform using platform-dependent tools.

This paper is based on the intuition that we can benefit that a code generator is often a member of a code generators family to automatically compare the performance of the different generated codes that comes from the same source program. Based on this comparison, we can automatically detect singular resource consumptions. As a result we proposes an approach to automatically compare and detect inconsistencies between code generators. This approach provides a runtime monitoring infrastructure based on system containers, as execution platforms, to compare the generated code performance. We evaluate our approach by analyzing the non-functional properties of HAXE code generator. Haxe is a popular high-level programming language[1] that involves a family of cross-platform code generators able to generate code to different targeted platforms. Our experimental results show that our approach is able to detect performance inconsistencies within HAXE code generators.

In this paper, we make the following contributions:

- We propose a fully automated micro-service infrastructure to ensure the deployment and monitoring of generated code. This paper focuses on the relationship between runtime execution of generated code and resource consumption profiles (memory usage).

- We also report the results of an empirical study by evaluating the non-functional properties of HAXE code generators. The obtained results provide evidence to support the claim that our proposed infrastructure is effective.

The paper is organized as follows. Section II describes the motivations behind this work by discussing three examples of code generator families. Section III presents an overview of our approach to automatically perform non-functional tests of code generator families. In particular, we present our infrastructure for non-functional testing of code generators using micro-services. The evaluation and results of our experiments are discussed in Section IV. Finally, related work, concluding remarks and future work are provided in Sections V and VI.

2. Motivations

2.1 Code generator families example

In different domain, the use of code generator is a common practice. We can cite three approaches that intensively develop and use code generators.

a. Haxe. Haxe is an open source toolkit for cross-platform development which compiles to a number of different programming platforms, including JavaScript, Flash, PHP, C++, C# and Java. Haxe involves many features: the Haxe language, multi-platform compilers, and different native libraries. The Haxe language is a high-level programming language which is strictly typed. This language supports both functional programming and object-oriented programming paradigms. It has a common type hierarchy, making certain API available on every targeted platform. Moreover, Haxe comes with a set of code generators that translate manually-written code (in Haxe language) to different target languages and platforms. Haxe code can be compiled for applications running on desktop, mobile and web platforms. Compilers ensure the correctness of user code in terms of syntax and type safety. Haxe comes also with a set of standard libraries that can be used on all supported targets and platform-specific libraries for each of them. One of the main usage of Haxe is to develop Cross-Platform Games or Cross-Platform libraries that can run on mobile, on the Web or on a Desktop. This project is popular (more than 1440 stars on github).

b. ThingML. ThingML is a modeling language for embedded and distributed systems. The idea of ThingML is to develop a practical model-driven software engineering toolchain which targets resource constrained embedded systems such as low-power sensor and microcontroller based devices. ThingML is developed as a domain-specific modeling language which includes concepts to describe both software components and communication protocols. The formalism used is a combination of architecture models, state machines and an imperative action language. The ThingML toolset provides a code generators family to translate ThingML to C, Java and JavaScript. It includes a set of variants for the C and JavaScript code generators to target different embedded system and their constraints. This project is still confidential but it is a good candidate to represent the modeling community practices.

c. TypeScript. TypeScript is a typed superset of JavaScript that compiles to plain JavaScript. In fact, it does not compile to only one version of JavaScript. It can transform typeScript to EcmaScript 3, 5 or 6. It can generate JavaScript that uses different system modules ('none', 'commonjs', 'amd', 'system', 'umd', 'es6', or 'es2015'). [2] This project is popular (more than 12,619 stars on github).

---

[1] 1442 github stars

[2] Each of this variation point can target different code generators (function *emitES6Module* vs *emitUMDModule* in emitter.ts for example).

2.2 Functional correctness of a family of code generators

A reliable and accepted way to increase confidence in the correct functioning of code generators is to validate and check the functionality of generated code, which is common practice for compiler validation and testing. Therefore, developers try to check the syntactic and semantic correctness of generated code by means of different techniques such as static analysis, test suites, etc., and ensure that the code is behaving correctly. In model-based testing for example [9, 19], testing code generators focuses on testing the generated code against its design. Thus, the model and the generated code are executed in parallel, by means of simulations, with the same set of test suites. Afterwards, the two outputs are compared with respect to certain acceptance criteria. Test cases, in this case, can be designed to maximize the code or model coverage [20].

Based on the three sample projects presented above, we can eventually observe on github that all of them use unit tests to check the correctness of the code generator and target language features.

2.3 Non-Functional correctness of a family of code generators

Code generators have to respect different requirements which preserve software reliability and quality [3]. A non-efficient code generator might generate defective software artifacts (code smells) that violates common software engineering practices. Thus, poor-quality code can affect system reliability and performance (e.g., high resource usage, low execution speed, etc.). Thus, another important aspect of code generator's testing is to test the non-functional properties of produced code. Proving that the generated code is functionally correct is not enough to claim the effectiveness of the code generator under test. In looking at the three motivating example, ThingML and TypeScript do not provide any specific test to check the consistency of the different memory usage or CPU consumption regarding the different code generators and their variants. Haxe provides two test cases [3] to benchmark the resulting generated code. One serves to benchmark an example in which object allocations are deliberately (over) used to measure how memory access/GC mixes with numeric processing in different target languages. The second test mainly bench the network speed for each target platform.

However, based on existing benchmarks between technical environments [8], in comparing the behavior of the generated code of a large data set of programs, our feeling is that we could detect inefficient code generators. The kinds of error that we track are the following:

- the lack of use of a **specific function that exists in the standard library** of the targeted language that can

---
[3] https://github.com/HaxeFoundation/haxe/tree/development/tests/benchs

speed or reduce the memory consumption of the resulting program.

- the lack of use of **a specific type that exists in the standard library** of the targeted language that can speed or reduce the memory consumption of the resulting program.

- the lack of use of **a specific language feature in a targeted language** that can speed or reduce the memory consumption of the resulting program.

The main difficulties is the fact that, for testing the non functional properties of a code generator, we cannot just observe the execution of the code generator but we have to observe and compare the execution of the generated program. Even if there is no exact oracle to detect inconsistencies, we could benefit from the family of code generators to compare the behavior of several programs generated from the same source that runs atop different technical stacks.

Next section discusses the common process used by a developer to automatically test the performance of a generated code and illustrate how we can benefit from the code generators families to identify singular behavior.

3. Approach overview

3.1 Non-Functional testing of a family of code generators: a common process

Figure 1 summarizes a general overview of the different steps that are involved together to ensure the code generation and non-functional testing of produced code from design time to run time. We distinguish 4 major steps: the software design using high-level system specifications, code generation by means of code generators, code execution, and non-functional testing of generated code.

In the first step, software developers have to define, at design time, software's behavior using a high-level abstract language (DSLs, models, program). Afterwards, developers can use platform-specific code generators to ease the software development and generate automatically code that targets different languages and platforms. We depict, in Figure 1, three code generators capable to generate code in three software programming languages (JAVA, C# and C++). In this step, code generators transform the previously designed model to produce, as a consequence, software artifacts for the target platform.

In the next step, generated software artifacts (e.g., JAVA, C#, C++, etc.) are compiled, deployed and executed across different target platforms (e.g., Android, ARM/Linux, JVM, x86/Linux, etc.). Thus, several code compilers are needed to transform source code to machine code (binaries) in order to get executed.

Finally, to perform the non-functional testing of generated code, developers have to collect, visualize and compare information about the performance and efficiency of running code across the different platforms. Therefore, they gener-

Figure 1. An overall overview of the different processes involved to ensure the code generation and non-functional testing of produced code from design time to run time: the classical way

ally use several platform-specific profilers, trackers, instrumenting and monitoring tools in order to find some inconsistencies or bugs during code execution [2, 5]. Ensuring the code quality of generated code can refer to several nonfunctional properties such as code size, resource or energy consumption, execution time, among others [12]. Finding inconsistencies within code generators involves analyzing and inspecting the code and that, for each execution platform. For example, one of the methods to handle that is to analyze the memory footprint of software execution and find memory leaks. Developers then, can inspect the source code and the generated code and find some parts of the code-base that have triggered this issue. Therefore, software testers generally use to report statistics about the performance of generated code in order to fix, refactor, and optimize the code generation process. Our approach aims to automate the three last steps: generate code for a set of platforms, execute code on top of different platforms, monitor and compare the execution.

3.2 An Infrastructure for Non-functional Testing Using System Containers

In general, there are many non-functional requirements that can be evaluated by software testers such as performance (execution time), code quality, robustness, resource usage, *etc*. In this paper, our approach focuses on the non-functional properties related to the generated code performance in terms of resource consumption (memory and CPU).

In fact, to assess the performance/non-functional properties of generated code many system configurations (i.e., execution environments) must be considered. Running different applications (i.e., generated code) with different configurations on one single machine is complex a single system has limited resources and this can lead to performance regressions. Moreover, each execution environment comes with a collection of appropriate tools such as compilers, code generators, debuggers, profilers, etc. Therefore, we need to de-

ploy the test harness, i. e.the produced binaries, on an elastic infrastructure that provides to compiler user facilities to ensure the deployment and monitoring of generated code in different environment settings.

Consequently, we rely on an infrastructure that must provide the following feature. The infrastructure must support to automatically:

1. Deploy the generated code, its dependencies and its execution environments

2. Execute the produces binaries in an isolated environment

3. Monitor the execution

4. Gather performance metrics (CPU, Memory, etc.)

To get this four main feature, the testing infrastructure is based on system containers [15]. First, instead of configuring all code generators under test (GUTs) within the same host machine, we wrap each GUT within a system container. Afterwards, a new instance of the container is created to enable the execution of generated code in a isolated monitored and configured environment. Meanwhile, we start our runtime testing components. A monitoring component collects usage statistics of all running containers and save them at runtime in the time series database component. Thus, we can compare later informations about the resource usage of generative program and detect inconsistencies in the generated code.

The rest of this section details the technical choices we made to synthesis a testing infrastructure to detect inconsistencies between code-generator regarding memory and CPU consumption.

3.3 Technical implementation

The general overview of the technical implementation is shown in Figure 2. In the following subsections, we describe the deployment and testing architecture of generated code within system containers.
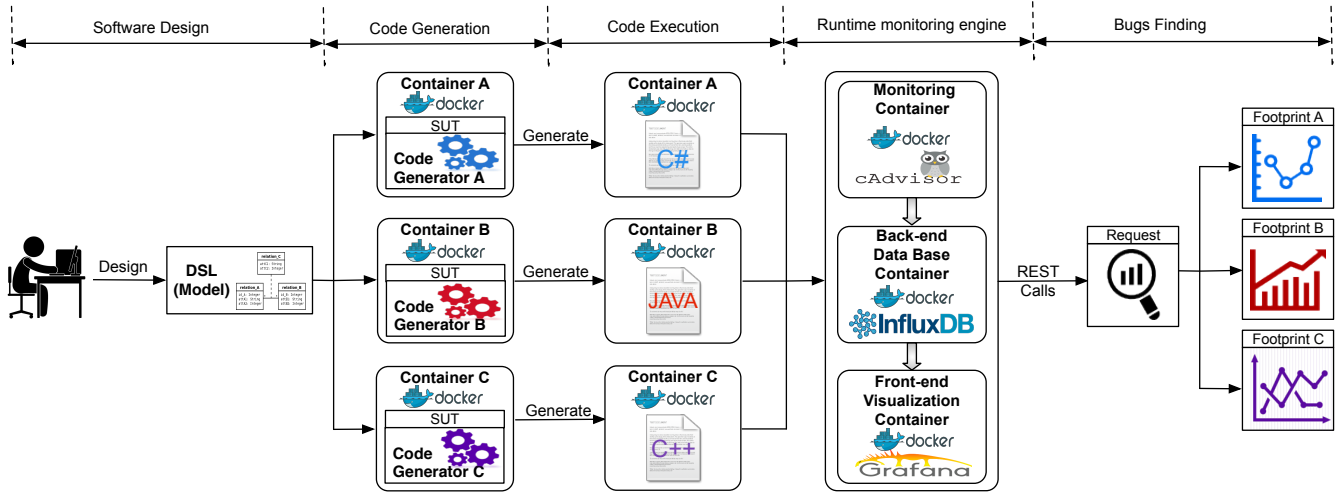
Figure 2. A technical overview of the different processes involved to ensure the code generation and non-functional testing of produced code from design time to runtime.

3.3.1 System Containers as Execution platforms

Before starting to monitor and test applications, we have to deploy generated code on different components to ease containers provisioning and profiling. We aim to use Docker Linux containers to monitor the execution of different generated artifacts in terms of resource usage [10]. Docker[4] is an open source container engine that automates the deployment of any application as a lightweight, portable, and self-sufficient container that runs virtually on a host machine. Using Docker, we can define pre-configured applications and servers to host as virtual images. We can also define the way the service should be deployed in the host machine using configuration files called Docker files. We use the public Docker registry[5] for saving, and managing all our Docker images. We can then instantiate different containers from these Docker images.

Therefore, each generated code is executed individually inside an isolated Linux container. By doing so, we ensure that each executed program runs in isolation without being affected by the host machine or any other processes. Moreover, since a container is cheap to create, we are able to create too many containers as long as we have new programs to execute. Since each program execution requires a new container to be created, it is crucial to remove and kill containers that have finished their job to eliminate the load on the system. We run the experiment on top of a private datacenter that provide a bare-metal installation of docker and docker swarm. On a single machine, containers/softwares are running sequentially and we pin p cores and n Gbytes of memory for each container[6]. Once the execution is done, resources reserved for the container are automatically released

to enable spawning next containers. Therefore, the host machine will not suffer too much from performance trade-offs.

In short, the main advantages of this approach are:

- The use of containers induces less performance overhead and resource isolation compared to using a full stack virtualization solution [16]. Indeed, instrumentation and monitoring tools for memory profiling like Valgrind [11] can induce too much overhead.

- Thanks to the use of Dockerfiles, the proposed framework can be configured by software testers in order to define the code generators under test (e. g., code generator version, dependencies, etc.), the host IP and OS, the DSL design, the optimization options, etc. Thus, we can use the same configured Docker image to execute different instances of generated code. For hardware architecture, containers share the same platform architecture as the host machine (e.g., x86, x64, ARM, etc.).

- Docker uses Linux control groups (Cgroups) to group processes running in the container. This allows us to manage the resources of a group of processes, which is very valuable. This approach increases the flexibility when we want to manage resources, since we can manage every group individually. For example, if we would evaluate the non-functional requirements of generated code within a resource-constraint environment, we can request and limit resources within the execution container according to the needs.

- Although containers run in isolation, they can share data with the host machine and other running containers. Thus, non-functional data relative to resource consumption can be gathered and managed by other containers (i. e., for storage purpose, visualization)

---

[4] https://www.docker.com

[5] https://hub.docker.com/

[6] p and n can be cofigured

3.3.2 Runtime Testing Components

In order to test our running applications within Docker containers, we aim to use a set of Docker components to ease the extraction of resource usage information.

Monitoring Component This container provides an understanding of the resource usage and performance characteristics of our running containers. Generally, Docker containers rely on Cgroups file systems to expose a lot of metrics about accumulated CPU cycles, memory, block I/O usage, etc. Therefore, our monitoring component automates the extraction of runtime performance metrics stored in Cgroups files. For example, we access live resource consumption of each container available at the Cgroups file system via stats found in *"/sys/fs/cgroup/cpu/docker/(longid)/"* (for CPU consumption) and *"/sys/fs/cgroup/memory/docker/(longid)/"* (for stats related to memory consumption). This component will automate the process of service discovery and metrics aggregation for each new container. Thus, instead of gathering manually metrics located in Cgroups file systems, it extracts automatically the runtime resource usage statistics relative to the running component (i.e., the generated code that is running within a container). We note that resource usage information is collected in raw data. This process may induce a little overhead because it does very fine-grained accounting of resource usage on running container. Fortunately, this may not affect the gathered performance values since we run only one version of generated code within each container. To ease the monitoring process, we integrate cAdvisor, a Container Advisor[7]. cAdvisor monitors service containers at runtime.

However, cAdvisor monitors and aggregates live data over only 60 seconds interval. Therefore, we record all data over time since container's creation in a time-series database. It allows the code-generator testers to run queries and define non-functional metrics from historical data. Thereby, to make gathered data truly valuable for resource usage monitoring, we link our monitoring component to a back-end database component.

Back-end Database Component This component represents a time-series database back-end. It is plugged with the previously described monitoring component to save the non-functional data for long-term retention, analytics and visualization.

During the execution of generated code, resource usage stats are continuously sent to this component. When a container is killed, we are able to access to its relative resource usage metrics through the database. We choose a time series database because we are collecting time series data that correspond to the resource utilization profiles of programs execution.

We use InfluxDB[8], an open source distributed time-series database as a back-end to record data. InfluxDB allows the user to execute SQL-like queries on the database. For example, the following query reports the maximum memory usage of container "$generated\_code\_v1$" since its creation:

```sql
select max (memory_usage) from stats
where container_name='generated_code_v1'
```

To give an idea about the data gathered from the monitoring component and stored in the time-series database, we describe in Table 1 these collected metrics:

Metric	Description
Name	Container Name
T	Elapsed time since container's creation
Network	Stats for network bytes and packets in an out of the container
Disk IO	Disk I/O stats
Memory	Memory usage
CPU	CPU usage

Table 1. Resource usage metrics recorded in InfluxDB

Apart from that, our framework provides also information about the size of generated binaries and the compilation time needed to produced code. For instance, resource usage statistics are collected and stored using these two components. It is relevant to show resource usage profiles of running programs overtime. To do so, we present a front-end visualization component for performance profiling.

Front-end Visualization Component Once we gather and store resource usage data, the next step is visualizing them. That is the role of the visualization component. It will be the endpoint component that we use to visualize the recorded data. Therefore, we provide a dashboard to run queries and view different profiles of resource consumption of running components through web UI. Thereby, we can compare visually the profiles of resource consumption among containers. Moreover, we use this component to export the data currently being viewed into static CSV document. So, we can perform statistical analysis on this data to detect inconsistencies or performance anomalies. As a visualization component, we use Grafana[9], a time-series visualization tool available for Docker.

3.3.3 Technical Architecture Overview

To summarize, we present, in Figure 2, a technical overview of the different components involved in our approach to perform the non-functional testing of code generators.

---

[7] https://github.com/google/cadvisor

[8] https://github.com/influxdata/influxdb

[9] https://github.com/grafana/grafana

First, instead of configuring all code generators under test (GUTs) within the same host machine (as we presented in the motivation section), our tool wrap each GUT within a container. To do so, we create a new configuration image for each GUT (i.e., the Docker image) where we install all the libraries, compilers, and dependencies needed to ensure the code generation and compilation. Thereby, the GUT produce code within multiple instances of preconfigured Docker images. Afterwards, a new instance of the container is created to enable the execution of generated code in a isolated and configured environment. In this step, new configurations are needed to run each job without any issue (i.e., define the execution environment within the container, the host machine, the target platform, limit the resources, etc.). For our case, a job represents a generated and compiled program using one of the GUTs. Meanwhile, we start our runtime testing components (e.g., cAdvisor, InfluxDB and Grafana). The monitoring component collects usage statistics of all running containers and save them at runtime in the time series database component. The visualization component comes later to allow end users to define performance metrics and draw up charts. The use of the front-end visualization component is optional and we can directly access to information stored in the database through REST API calls. Thus, we can observe later informations about the resource usage of generated code and detect inconsistencies in the generated code by GUTs.

4. Evaluation

So far, we have presented a sound procedure and automated component-based framework for extracting the non-functional properties of generated code. In this section, we evaluate the implementation of our approach by explaining the design of our empirical study and the different methods we used to assess the effectiveness of our approach. The experimental material is available for replication purposes[10].

4.1 Experimental Setup

4.1.1 Code Generators Under Test: HAXE compilers

In order to test the applicability of our approach, we conduct experiments on a popular high-level programming language called HAXE and its code generators.

Haxe comes with a set of compilers that translate manually-written code (in Haxe language) to different target languages and platforms.

The process of code transformation and generation can be described as following: Haxe compilers analyzes the source code written in Haxe language then, the code is checked and parsed into a typed structure, resulting in a typed abstract syntax tree (AST). This AST is optimized and transformed afterwards to produce source code for target platform/language.

Haxe offers the option of choosing which platform to target for each program using a command-line tool. Moreover, some optimizations and debugging information can be enabled through CLI but in our experiments, we did not turned on any further options.

4.1.2 Cross-platform Benchmark

One way to prove the effectiveness of our approach is to create benchmarks. Thus, we use the Haxe language and its code generators to build a cross-platform benchmark. The proposed benchmark is composed of a collection of cross-platform libraries that can be compiled to different targets. In these experiments, we consider five Haxe code generators to test: Java, JS, C++, CS, and PHP code generator. To select a subset of existing haxe library that can be used in out experiment, we crawl github and we use the haxe library repository[11]. We select the seven libraries that has the best code coverage score.

In fact, each Haxe library comes with an API and a set of test suites. These tests, written in Haxe, represent a set of unit tests that covers the different functions of the API. The main task of these tests is to check the correct functional behavior of generated programs once generated code is executed within the target platform. To prepare our benchmark, we have removed all the tests that fail to compile to our five targets (i.e., errors, crashes and failures) and we kept only test suites that are functionally correct to focus on performance. Moreover, we added manually new test cases to some libraries in order to extend the number of test suites. The number of test suites depends on the number of functions within the Haxe library.

We use then, these test suites then, to generate a load and stress the target library. This can be useful to study the impact of this load on the resource usage of the system. For example, if one test suite consumes a lot of resources for a specific target, then this could be explained by the fact that the code generator has produced code that is very greedy in terms of resources.

Thus, we run each test suite 1000 times to get comparable values in terms of resource usage. Table 2 describes the Haxe libraries that we have selected from this benchmark to evaluate our approach.

4.1.3 Evaluation Metrics Used

We use to evaluate the efficiency of generated code using the following non-functional metrics:

-Memory Usage (MU): It corresponds to the maximum memory consumption of the running container under test. Memory usage is measured in bytes.

*-Execution Time (T)*: Program execution time is measured in seconds.

Library	#TestSuites	Description
Color	19	Color conversion from/to any color space
Core	51	Provides extensions to many types
Hxmath	6	A 2D/3D math library
Format	4	Format library such as dates, number formats
Promise	3	Library for lightweight promises and futures
Culture	4	Localization library for Haxe
Math	3	Generation of random values

Table 2. Description of selected benchmark libraries

We recall that our tool is able to evaluate other non-functional properties of generated code such as code generation time, compilation time, code size, CPU usage.

4.1.4 Setting up Infrastructure
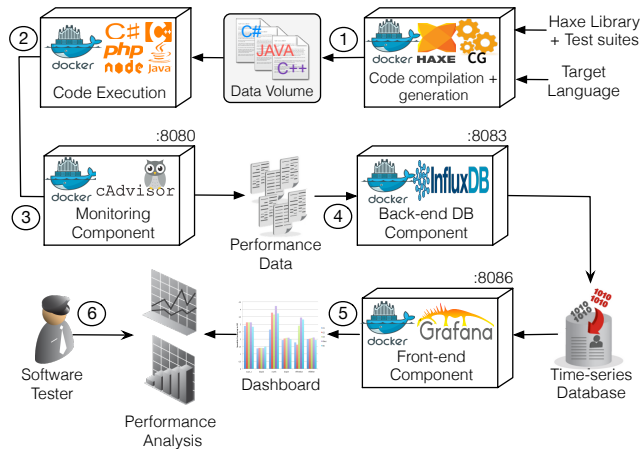
Figure 3. Infrastructure settings for running experiments

To assess our approach, we configure our previously proposed container-based infrastructure for non-functional testing of code generators in order to run experiments on the Haxe case study. Figure 3 shows a big picture of the testing and monitoring infrastructure considered in these experiments.

¡¡¡¡¡¡¡ HEAD ======= ¿¿¿¿¿¿¿ cbf285ae38b983232649770826a12d9970c12b41

First, we create a new Docker image in where we install the Haxe code generators and compilers (through the configuration file "Dockerfile"). Then a new instance of that image is created. It takes as an input the Haxe library we would to test and the list of test suites (step 1). It produces as an output the source code and binaries that have to be executed. These files are saved in a shared repository. In Docker environment, this repository is called Data Volume. A data volume is a specially-designated directory within containers that share data with the host machine. So, when we execute

Benchmark	TestSuite	Std_dev	TestSuite	Std_dev	TestSuite	Std_dev
Color	TS1	0,55	TS8	0,24	TS15	0,73
	TS2	0,29	TS9	0,22	TS16	0,12
	TS3	0,34	TS10	0,10	TS17	0,31
	TS4	2,51	TS11	0,17	TS18	0,34
	TS5	1,53	TS12	0,28	TS19	120,61
	TS6	43,50	TS13	0,33		
	TS7	0,50	TS14	1,88		
Core	TS1	0,35	TS18	0,16	TS35	1,30
	TS2	0,07	TS19	0,60	TS36	1,13
	TS3	0,30	TS20	5,79	TS37	2,02
	TS4	27299,89	TS21	0,47	TS38	0,26
	TS5	6,12	TS22	2,74	TS39	0,16
	TS6	21,90	TS23	2,14	TS40	8,12
	TS7	0,41	TS24	3,79	TS41	5,45
	TS8	0,28	TS25	0,19	TS42	0,11
	TS9	0,78	TS26	0,13	TS43	1,41
	TS10	1,82	TS27	5,59	TS44	1,56
	TS11	180,68	TS28	1,71	TS45	0,11
	TS12	185,02	TS29	0,26	TS46	1,04
	TS13	128,78	TS30	0,44	TS47	0,23
	TS14	0,71	TS31	1,71	TS48	1,34
	TS15	0,12	TS32	2,42	TS49	1,86
	TS16	0,65	TS33	8,29	TS50	1,28
	TS17	0,26	TS34	5,25	TS51	3,53
Hxmath	TS1	31,65	TS3	30,34	TS5	0,40
	TS2	4,27	TS4	0,25	TS6	0,87
Format	TS1	0,28	TS3	95,36	TS4	1,49
	TS2	64,94				
Promise	TS1	0,29	TS2	13,21	TS3	1,21
Culture	TS1	0,13	TS3	0,13	TS4	1,40
	TS2	0,10				
Math	TS1	642,85	TS2	28,32	TS3	24,40

¡¡¡¡¡¡¡ HEAD

Table 3. The comparison results of running each test suite across five target languages: the metric used is the standard deviation between execution times

=======

Table 4. Execution speed results
¿¿¿¿¿¿¿ cbf285ae38b983232649770826a12d9970c12b41

the generated test suites, we provide a shared volume with the host machine so that, binaries can be executed in the execution container (Step 2). In fact, for the code execution we created, as well, a new Docker image in where we install all execution tools and environments such as php interpreter, NodeJS, etc.

In the meantime, while running test suites inside the container, we collect runtime resource usage data using cAdvisor (step 3). The cAdvisor Docker image does not need any configuration on the host machine. We have just to run it on our host machine. It will then have access to resource usage and performance characteristics of all running containers. This image uses the Cgroups mechanism described previously to collect, aggregate, process, and export ephemeral real-time information about running containers. Then, it reports all statistics via web UI to view live resource consumption of each container. cAdvisor has been widely used in different projects such as Heapster[12] and Google Cloud Platform[13]. In this experiment, we choose to gather information about the memory usage of running container. Afterwards,

---

[12] https://github.com/kubernetes/heapster

[13] https://cloud.google.com/

Benchmark	TestSuite	Std_dev	TestSuite	Std_dev	TestSuite	Std_dev
Color	TS1	10,19	TS8	1,23	TS15	14,44
	TS2	1,17	TS9	1,95	TS16	1,13
	TS3	0,89	TS10	1,27	TS17	0,72
	TS4	30,34	TS11	0,57	TS18	0,97
	TS5	31,79	TS12	1,11	TS19	777,32
	TS6	593,05	TS13	0,46		
	TS7	12,14	TS14	45,90		
Core	TS1	1,40	TS18	1,00	TS35	14,13
	TS2	1,17	TS19	20,37	TS36	32,41
	TS3	0,60	TS20	128,23	TS37	22,72
	TS4	403,15	TS21	24,38	TS38	2,19
	TS5	41,95	TS22	76,24	TS39	0,26
	TS6	203,55	TS23	18,82	TS40	126,29
	TS7	19,69	TS24	72,01	TS41	31,01
	TS8	0,78	TS25	0,21	TS42	0,93
	TS9	30,41	TS26	2,30	TS43	50,36
	TS10	57,19	TS27	101,53	TS44	12,56
	TS11	68,92	TS28	43,67	TS45	0,91
	TS12	74,19	TS29	0,90	TS46	27,28
	TS13	263,99	TS30	4,02	TS47	1,10
	TS14	19,89	TS31	52,35	TS48	15,40
	TS15	0,30	TS32	134,75	TS49	37,01
	TS16	28,29	TS33	82,66	TS50	23,29
	TS17	1,16	TS34	89,57	TS51	1,28
Hxmath	TS1	444,18	TS3	425,65	TS5	17,69
	TS2	154,80	TS4	0,96	TS6	46,13
Format	TS1	0,74	TS3	255,36	TS4	8,40
	TS2	106,87				
Promise	TS1	0,30	TS2	58,76	TS3	20,04
Culture	TS1	1,28	TS3	0,58	TS4	15,69
	TS2	4,51				
Math	TS1	1041,53	TS2	234,93	TS3	281,12

¡¡¡¡¡¡ HEAD

Table 5. The comparison results of running each test suite across five target languages: the metric used is the standard deviation between memory consumptions

=======

Table 6. Memory consumption speed results
¿¿¿¿¿¿ cbf285ae38b983232649770826a12d9970c12b41

we record these data into a new time-series database using our InfluxDB back-end container (step 4).

Next, we run Grafana and we link it to InfluxDB. Grafana can request data from the database. We recall that InfluxDB also provides a web UI to query the database and show graphs (step 5). But, Grafana let us display live results over time in much pretty looking graphs. Same as InfluxDB, we use SQL queries to extract non-functional metrics from the database for visualization and analysis (step 6). In our experiment, we are gathering the maximum memory usage values without presenting the graphs of resource usage profiles.

To obtain comparable and reproducible results, we use the same hardware across all experiments: a farm of AMD A10-7700K APU Radeon(TM) R7 Graphics processor with 4 CPU cores (2.0 GHz), running Linux with a 64 bit kernel and 16 GB of system memory. We reserve one core and 4Gbytes of memory for each running container.

4.2 Experimental Results

4.2.1 Evaluation using the standard deviation

We now conduct experiments based on the Haxe benchmark. We run each test suite 1K times and we report the execution time and memory usage across the different target lan-

guages: Java, JS, C++, CS, and PHP. The goal of running these experiments is to observe and compare the behavior of generated code regarding the testing load. We recall, as mentioned in the motivation, that we are not using any oracle function to detect inconsistencies. However, we rely on the comparison results across different targets to define code generator inconsistencies. Thus, we use, as a quality metric, the standard deviation to quantify the amount of variation among execution traces (i.e., memory usage or execution time) and that for the five target languages. A low standard deviation of a test suite execution, indicates that the data points (execution time or memory usage data) tend to be close to the mean which we consider as an acceptable behavior. On the other hand, a high standard deviation indicates that one or more data points are spread out over a wider range of values which can be more likely interpreted as a code generator inconsistency.

In Table 3, we report the comparison results of running the benchmark in terms of execution speed. At the first glance, we can clearly see that all standard deviations are more mostly close to 0 - 8 interval. It is completely normal to get such small deviations, because we are comparing the execution time of test suites that are written in heterogeneous languages and executed using different technologies (e.g., interpreters for PHP, JVM for JAVA, etc.). So, it is expected to get a small deviation between the execution times after running the test suite in different languages. However, we remark in the same table, that there are some variation points where the deviation is relatively high. We count 8 test suites where the deviation is higher than 60 (highlighted in gray). We choose this value (i.e., standard deviation = 60) as a threshold to designate the points where the variation is extremely high. Thus, we consider values higher than 60 as a potential possibility that a non-functional bug can occur. These variations can be explained by the fact that the execution speed of one or more test suites varies considerably from one language to another. This argues the idea that the code generator has produced a suspect behavior of code for one or more target language. We provide later better explanation in order to detect the faulty code generators.

Similarly, table 4 resumes the comparison results of test suites execution regarding memory usage. The variation in this experiment are more important than previous results. This can be argued by the fact that the memory utilization and allocation patterns are different for each language. Nevertheless, we can recognize some points where the variation is extremely high. Thus, we choose a threshold value equal to 400 and we highlighted, in gray, these extreme points. Thus, we detected 6 test suites where the the variation is extremely high. One of the reasons that caused this variation may occur when the test suite executes some parts of the code (in a specific language) that are so greedy in terms of resources. This may be not the case when the variation is lower than 10 for example. We assume then that the faulty

	JS		JAVA		C++		CS		PHP	
	Time	**Factor**	**Time**	**Factor**	**Time**	**Factor**	**Time**	**Factor**	**Time**	**Factor**
Color_TS19	4,52	x1.0	8,61	x1,9	10,73	x2,4	14,99	x3,3	279,27	x61,8
Core_TS4	665,78	x1.0	416,85	x0,6	699,11	x1,1	1161,29	x1,7	61777,21	x92,8
Core_TS11	4,27	x1.0	1,80	x0,4	1,57	x0,4	5,71	x1,3	407,33	x95,4
Core_TS12	4,71	x1.0	2,06	x0,4	1,60	x0,3	5,36	x1,1	417,14	x88,6
Core_TS13	6,26	x1.0	5,91	x0,9	11,04	x1,8	14,14	x2,3	297,21	x47,5
Format_TS2	2,31	x1.0	2,10	x0,9	1,81	x0,8	6,08	x2,6	148,24	x64,1
Format_TS3	5,40	x1.0	5,03	x0,9	7,67	x1,4	12,38	x2,3	220,76	x40,9
Math_TS1	3,01	x1.0	12,51	x4,2	16,30	x5,4	14,14	x4,7	1448,90	x481,7

Table 7. Raw data values of test suites that led to the highest variation in terms of execution time

	JS		JAVA		C++		CS		PHP	
	Memory	**Factor**	**Memory**	**Factor**	**Memory**	**Factor**	**Memory**	**Factor**	**Memory**	**Factor**
Color_TS6	900,70	x1.0	1362,55	x1,5	2275,49	x2,5	1283,31	x1,4	758,79	x0,8
Color_TS19	253,01	x1.0	819,92	x3,2	923,99	x3,7	327,61	x1,3	2189,86	x8,7
Core_TS4	303,09	x1.0	768,22	x2,5	618,42	x2	235,75	x0,8	1237,15	x4,1
Hxmath_TS1	104,00	x1.0	335,50	x3,2	296,43	x2,9	156,41	x1,5	1192,98	x11,5
Hxmath_TS3	111,68	x1.0	389,73	x3,5	273,12	x2,4	136,49	x1,2	1146,05	x10,3
Math_TS1	493,66	x1.0	831,44	x1,7	1492,97	x3	806,33	x1,6	3088,15	x6,3

Table 8. Raw data values of test suites that led to the highest variation in terms of memory usage

	JS		JAVA		C++		CS		PHP	
	Time	**Factor**	**Time**	**Factor**	**Time**	**Factor**	**Time**	**Factor**	**Time**	**Factor**
Color_TS19	4,52	x1.0	8,61	x1,9	10,73	x2,4	14,99	x3,3	279,27	x61,8
Core_TS4	665,78	x1.0	416,85	x0,6	699,11	x1,1	1161,29	x1,7	61777,21	x92,8
Core_TS11	4,27	x1.0	1,80	x0,4	1,57	x0,4	5,71	x1,3	407,33	x95,4
Core_TS12	4,71	x1.0	2,06	x0,4	1,60	x0,3	5,36	x1,1	417,14	x88,6
Core_TS13	6,26	x1.0	5,91	x0,9	11,04	x1,8	14,14	x2,3	297,21	x47,5
Format_TS2	2,31	x1.0	2,10	x0,9	1,81	x0,8	6,08	x2,6	148,24	x64,1
Format_TS3	5,40	x1.0	5,03	x0,9	7,67	x1,4	12,38	x2,3	220,76	x40,9
Math_TS1	3,01	x1.0	12,51	x4,2	16,30	x5,4	14,14	x4,7	1448,90	x481,7

Table 9. My caption

code generator, in this case, represents a threat for software quality since it can generate a code that is very resource consuming.

The inconsistencies we are trying to find here are more related to the incorrect memory utilization patterns produced by the faulty code generator. Such inconsistencies may come from an inadequate type usage, high resource instantiation, etc.

4.2.2 Analysis

Now that we have observed the non-functional behavior of test suites execution in different languages, we can analyze the extreme points we have detected in previous tables to dig more in deep the source of such deviation. For that reason, we present in table 5 and 6 the raw data values of these extreme test suites in terms of execution time and memory usage.

Table 5 is showing the execution time of each test suite in a specific target language. We provides also factors of execution times among test suites running in different languages by taking as a baseline the JS version. Prior to this, wa calculated the average execution time and memory usage for all test suites per language (for all benchmark libraries). We found that JS has the lowest memory usage and execution time therefore, we choose it as a baseline. Based on these results, we can clearly see that the PHP code has the lowest performance with a factor ranging from x40.9 for testsuite 3 in benchamrk Format (Format_TS3) to x481.7 for Math_TS1. We remark also that running Core_TS4 takes 61777 seconds (almost 17 hours) compared to a 416 seconds (around 6 minutes) in JAVA which is a very large gap. The highest factor detected for other languages ranges from x0.3 to x5.4 which is not negligible but it represents a small deviation compared to PHP version. While it is true that we are comparing different versions of generated code, it was expected to get some variations while running test cases in terms of execution time. However, in the case of PHP code generator it is far to be a simple variation but it is more likely to be a code generator inconsistency that led to such performance regression.

Meanwhile, we gathered information about the points that led to the highest standard derivation in terms of memory usage. Table 6 shows these results. We take as well the JS version as a baseline since it requires less memory. Again, PHP test suites consumed almost the highest amount of memory (five in six of test suites). However, for Color_TS6, C# version consumes the highest memory (x2.5 more than JS). For other test suites versions, the factor varies from x0.8 to x3.7. Besides the performance issues of PHP code generator presented in table 5, the results of memory usage confirm our claim since the PHP code has the highest memory utilization.

4.3 Discussions

Through these conducted experiments, we reached interesting results, some of which were unexpected. For instance, the poor performance of the Haxe-produced PHP code was certainly a surprise. The PHP target had a very large execution time and memory usage for almost every test suite execution. That is to say that PHP code generator is clearly generating non-efficient code regarding the non-functional properties. It is even possible to say that other code generators are not extremely efficient since we found that C++ code consumed, during one test suite execution (Color_TS6), more memory than PHP. But, we cannot say for sure that C++ code generator is buggy. Thus, we cannot make any assumption. Nevertheless, the only point which, at this stage, can be statistically made is that PHP code generator has a real performance issue that has to be fixed by code generators developers.

In attempting to understand the reasons of this PHP performance issues, we tried to take a look at the source code generated in PHP. As an example, we looked at the source code of Core_TS4. In fact, throughout the inspection of generated source code, we remark the intensive use of *"arrays"* in most of the functions under test. In most languages, arrays are fixed-sized, but this is not the case in PHP since they are allocated dynamically. The dynamic allocation of arrays leads to a slower write time because the memory locations needed to hold the new data is not already allocated. Thus, slow writing speed damages the performance of PHP code and impact the memory usage. This observation clearly confirm our early findings. The solution for this problem may be to use another type of object from the Standard PHP Library. As an alternative, *"SplFixedArray"* pre-allocates the necessary memory and allows a faster array implementation, thereby solving the issue of slower write times. [ref to spl benchs]

In short, the lack of use of specific types in native PHP standard library by the PHP code generator such as *SplFixedArray* shows a real impact on the non-functional behavior of generated code. In contrast, selecting carefully the adequate types and functions to generate code by code generators can lead to performance improvement.

5. Related Work

//The idea of employing container-based approach for testing code generators can be found in several publications.

Previous work on non-functional testing of code generators focuses on comparing, as oracle, the non-functional properties of hand-written code to automatically generated code [14, 17]. As an example, Strekelj et al. [18] implemented a simple 2D game in both the Haxe programming language and the native environment and evaluated the difference in performance between the two versions of code. They showed that the generated code through Haxe has better performance than hand-written code.

Cross-platform mobile development has been also part of the non-functional testing goals since many code generators are increasingly used in industry for automatic cross-platform development. in [7, 13], authors compare the performance of a set of cross-platform code generators and presented the most efficient tools.

Most of the previous work on code generators testing focuses on checking the correct functional behavior of generated code. Stuermer et al. [19] present a systematic test approach for model-based code generators. They investigate the impact of optimization rules for model-based code generation by comparing the output of the code execution with the output of the model execution. If these outputs are equivalent, it is assumed that the code generator works as expected. They evaluate the effectiveness of this approach by means of testing optimizations performed by the TargetLink code generator. They have used Simulink as a simulation environment of models.

In [9], authors presented a testing approach of the Genesys code generator framework which tests the translation performed by a code generator from a semantic perspective rather than just checking for syntactic correctness of the generation result.

In our approach, we provide a component-based infrastructure to compare non-functional properties of generated code rather than functional ones.

6. Conclusion and Future Work

In this paper we have described a new approach for testing and monitoring code generators families using a container-based infrastructure. We used a set of micro-services in order to provide a fine-grained understanding of resource consumption. To validate the approach, we use the proposed approach on an widely used code generator families: Haxe. This evaluation shows that we could find real issue in the existing code generators. In particular, we show that we could find two kinds of errors: the lack of use of a specific function and abstract type that exist in the standard library of a targeted language that can reduce the memory/CPU consumption of the resulting program.

As a current work, we currently discuss with the Haxe community to submit a patch with the first discoveries. We

are also conducting the same evaluation for ThingML and TypeScript. As a future work, we would like to better understand if it exists a threshold that provides a best precision for detecting performance issue in code generators.

Acknowledgments

References

[1] M. Brambilla, J. Cabot, and M. Wimmer. Model-driven software engineering in practice. *Synthesis Lectures on Software Engineering*, 1(1):1–182, 2012.

[2] N. Delgado, A. Q. Gates, and S. Roach. A taxonomy and catalog of runtime software-fault monitoring tools. *IEEE Transactions on software Engineering*, 30(12):859–872, 2004.

[3] M. Demertzi, M. Annavaram, and M. Hall. Analyzing the effects of compiler optimizations on application reliability. In *Workload Characterization (IISWC), 2011 IEEE International Symposium on*, pages 184–193. IEEE, 2011.

[4] R. France and B. Rumpe. Model-driven development of complex software: A research roadmap. In *2007 Future of Software Engineering*, pages 37–54. IEEE Computer Society, 2007.

[5] V. Guana and E. Stroulia. Chaintracker, a model-transformation trace analysis tool for code-generation environments. In *ICMT*, pages 146–153. Springer, 2014.

[6] V. Guana and E. Stroulia. How do developers solve software-engineering tasks on model-based code generators? an empirical study design. In *First International Workshop on Human Factors in Modeling (HuFaMo 2015). CEUR-WS*, pages 33–38, 2015.

[7] G. Hartmann, G. Stead, and A. DeGani. Cross-platform mobile development. *Mobile Learning Environment, Cambridge*, pages 1–18, 2011.

[8] R. Hundt. Loop recognition in c++/java/go/scala. 2011.

[9] S. Jörges and B. Steffen. Back-to-back testing of model-based code generators. In *International Symposium On Leveraging Applications of Formal Methods, Verification and Validation*, pages 425–444. Springer, 2014.

[10] D. Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux Journal*, 2014(239):2, 2014.

[11] N. Nethercote and J. Seward. Valgrind: a framework for heavyweight dynamic binary instrumentation. In *ACM Sigplan notices*, volume 42, pages 89–100. ACM, 2007.

[12] Z. Pan and R. Eigenmann. Fast and effective orchestration of compiler optimizations for automatic performance tuning. In *International Symposium on Code Generation and Optimization (CGO'06)*, pages 12–pp. IEEE, 2006.

[13] A. Pazirandeh and E. Vorobyeva. Evaluation of cross-platform tools for mobile development. 2015.

[14] J. Richard-Foy, O. Barais, and J.-M. Jézéquel. Efficient high-level abstractions for web programming. In *ACM SIGPLAN Notices*, volume 49, pages 53–60. ACM, 2013.

[15] S. Soltesz, H. Pötzl, M. E. Fiuczynski, A. Bavier, and L. Peterson. Container-based operating system virtualization: a scalable, high-performance alternative to hypervisors. In *ACM SIGOPS Operating Systems Review*, volume 41, pages 275–287. ACM, 2007.

[16] C. C. Spoiala, A. Calinciuc, C. O. Turcu, and C. Filote. Performance comparison of a webrtc server on docker versus virtual machine. In *2016 International Conference on Development and Application Systems (DAS)*, pages 295–298. IEEE, 2016.

[17] S. Stepasyuk and Y. Paunov. Evaluating the haxe programming language-performance comparison between haxe and platform-specific languages. 2015.

[18] D. Štrekelj, H. Leventić, and I. Galić. Performance overhead of haxe programming language for cross-platform game development. *International Journal of Electrical and Computer Engineering Systems*, 6(1):9–13, 2015.

[19] I. Stuermer, M. Conrad, H. Doerr, and P. Pepper. Systematic testing of model-based code generators. *IEEE Transactions on Software Engineering*, 33(9):622, 2007.

[20] I. Stürmer, D. Weinberg, and M. Conrad. Overview of existing safeguarding techniques for automatically generated code. In *ACM SIGSOFT Software Engineering Notes*, volume 30, pages 1–6. ACM, 2005.